

# Representing Data

Dov Kruger

Department of Electrical and Computer Engineering  
Rutgers University

October 6, 2025



- This module describes how to represent data as bits
- By the end you will be able to
  - Represent unsigned and signed Integers
  - Convert binary to/from decimal
  - Convert binary to/from hexadecimal
  - Compute twos complement numbers
  - Identify when integers overflow
  - Identify how text is stored on computers (ASCII and Unicode)
  - Fixed point fractions
  - Floating point numbers



- With binary data, it's all context
- It is not obvious what bits represent
- In this module you will learn how the computer represents
  - Unsigned integers
  - Signed integers
  - Fractions
  - text



# Bits and Bytes

- Bit: Smallest unit of data (0,1)
- Byte: Group of 8 bits
- Type of data completely depends on context



# Combinations

- 2 bits: 00, 01, 10, 11
- 3 bits: 000, 001, 010, 011, 100, 101, 110, 111
- $n$  bits =  $2^n$  combinations
- 8 bits =  $2^8 = 256$  combinations
- 16 bits =  $2^{16} = 65536$  combinations
- 32 bits =  $2^{32} = 4294967296$  combinations



# Different Binary Data Representations

- Unsigned Integers
- Signed Integers
- ASCII text
- Unicode
- Fixed Point Fractions
- Floating Point



# Integer Representations

- Integers come in different sizes (8, 16, 32, 64 bits)
- Signed/Unsigned
- Little-endian/Big-endian

## Example

Big-endian	00000000 00000001	00000000	00000000	1
Little-endian	00000001 00000000	00000000	00000000	1



# Representing 3 bit Integers

bits	unsigned	signed
000	0	0
001	1	1
010	2	2
011	3	3
100	4	-4
101	5	-3
110	6	-2
111	7	-1

What happens if we add 1 to 7?

What happens if we subtract 1 from 0?





# Unsigned Integers

- In decimal, each digit represents a power of 10
- Digits are 0-9
- Example:  $123 = 1 * 10^2 + 2 * 10^1 + 3 * 10^0 = 100 + 20 + 3$
- In binary, each digit represents a power of 2
- Digits are 0-1
- Example:  $101 = 1 * 2^2 + 0 * 2^1 + 1 * 2^0 = 4 + 0 + 1 = 5$
- Example:  
 $101101 = 1 * 2^5 + 0 * 2^4 + 1 * 2^3 + 1 * 2^2 + 0 * 2^1 + 1 * 2^0 = 32 + 8 + 4 + 0 + 1 = 45$



# Practice

binary	decimal
00001100	
00101010	
01100000	
10110001	
11111111	



A number in a base  $b$  is represented as a sum of powers of  $b$

- Base 10 (digits 0-9):  $123 = 1 * 10^2 + 2 * 10^1 + 3 * 10^0$
- Base 2 (digits 0-1):  
 $1111111 = 1 * 2^6 + 1 * 2^5 + 1 * 2^4 + 1 * 2^3 + 1 * 2^2 + 1 * 2^1 + 1 * 2^0$
- Base 8 (digits 0-7):  
 $123 = 1 * 8^2 + 2 * 8^1 + 3 * 8^0 = 64 + 16 + 3 = 83$
- Base 16 (digits 0-9, A-F):  $123 = 7 * 16^1 + 11 * 16^0$



# Examples of Base numbers

- What is  $(543)_8$  in decimal?
- What is  $(101101)_2$  in decimal?
- What is  $(3AB)_{16}$  in decimal?
- 78651 cannot be base 8, why not?



## Base 8: Octal

Every digit in octal represents 3 bits

Oct	bits	—	Oct	bits
0	000		4	100
1	001		5	101
2	010		6	110
3	011		7	111

$$(123)_8 = 1 * 8^2 + 2 * 8^1 + 3 * 8^0 = 64 + 16 + 3 = 83$$

$$(456)_8 = 100101110$$



# Base 16: Hexadecimal

Hex	bits	—	Hex	bits
0	0000		8	1000
1	0001		9	1001
2	0010		A	1010
3	0011		B	1011
4	0100		C	1100
5	0101		D	1101
6	0110		E	1110
7	0111		F	1111



# Encoding a byte in Hexadecimal

- $D9 = 11011001$
- $AF = 10101111$
- $8C = 10001100$



# Converting Between Binary and Decimal

$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$
128	64	32	16	8	4	2	1

- Method 1: sum powers of 2
- $10010010 = 128 + 16 + 2 = 146$
- Method 2: start from left
  - Start with 1
  - For each digit, multiply by 2
  - If the digit is 1, add 1
  - Example:  $(((((1 * 2) * 2) * 2) + 1) * 2 * 2 * 2 + 1) * 2 = 146$





# Twos Complement Arithmetic

- Consider just 8 bit number for simplicity
- First bit is sign 0 = positive, 1 = negative
- To negate a number
  - invert all bits
  - add 1
  - the resulting number is the negative of the original
  - Example:  $5 = 00000101 \rightarrow 11111010 \rightarrow 11111011 = -5$
  - Example:  $17 = 00010001 \rightarrow 11101110 \rightarrow 11101111 = -17$
  - Example:  $-11 = 11110101 \rightarrow 00001010 \rightarrow 00001011 = 11$
  - Example:  
 $-128 = 10000000 \rightarrow 01111111 \rightarrow 00000000 = -128$



# Overflow and Underflow

- Overflow is when the result of a computation is too large to fit
- Underflow is the same in the negative direction
- Example: given 3-bit unsigned
  - $3 + 5 = 1000 = 8 = 000 = 0$
  - $4 + 6 = 1010 = 10 = 010 = 2$
  - $4 - 5 = 111 = 7$
- Example: given 3-bit signed
  - $3 + 2 = 101 = -2$
  - $2 - 3 = 111 = -1$
  - $3 + 1 = 100 = -4$



# Overflow and Underflow

- When a result is too large, store only the low  $n$  bits
- Example: 3 bits
  - $3 + 3 = 6$  (no overflow)
  - $4 + 4 = 8$  (too big)  $= 0$  (overflow)
  - $3 - 2 = 1$  (no overflow)
  - $3 - 4 = -1 = 7$  (underflow)



# Integer Data Types

bits		minval	maxval
8	signed	-128	127
8	unsigned	0	255
16	signed	-32768	32767
16	unsigned	0	65535
32	signed	-2147483648	2147483647
32	unsigned	0	4294967295
64	signed	-9223372036854775808	9223372036854775807
64	unsigned	0	18446744073709551615



# Integer Operations Overview

- Integer operations are fundamental in digital systems
- We'll cover arithmetic, logical, and shift operations
- Each operation will be demonstrated with 8-bit examples
- Understanding these operations is crucial for digital design



# Addition

- Addition is performed bit by bit, carrying over when necessary
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ + 00110110 \\ \hline 10011111 \end{array}$$

- Note: Overflow can occur if the result exceeds 8 bits



# Subtraction

- Subtraction is often implemented as addition with two's complement
- Example (8-bit):

$$\begin{array}{r} 1\ 11 \\ 01101001 \\ 00110110 \\ \hline 00010011 \end{array}$$

- Two's complement of 00110110 is 11001010



# Bitwise AND

- AND operation: 1 if both bits are 1, otherwise 0
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ \& 00110110 \\ \hline 00110000 \end{array}$$

- Often used for masking specific bits





# Bitwise OR

- OR operation: 1 if either bit is 1, otherwise 0
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ | 00110110 \\ \hline 01111111 \end{array}$$

- Useful for setting specific bits



# Bitwise XOR

- XOR operation: 1 if bits are different, 0 if same
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ \oplus 00110110 \\ \hline 01011111 \end{array}$$

- Often used in cryptography and error detection



# Bitwise NOT

- Inverts all bits: 1 becomes 0, 0 becomes 1
- Example (8-bit):

$$\begin{array}{r} \sim 01101001 \\ \hline 10010110 \end{array}$$



# Logical Left Shift

- Shifts all bits to the left, filling rightmost with 0
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \ll 2 \\ \hline 10100100 \end{array}$$

- Equivalent to multiplication by  $2^n$  for n shifts



# Logical Right Shift

- Shifts all bits to the right, filling leftmost with 0
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \ll 2 \\ \hline 10100100 \end{array}$$

- Equivalent to multiplication by  $2^n$  for n shifts (unsigned)



# Rotate Left

- All bits move to the left
- Bits that drop off the end come in on right
- Example (8-bit, rotate by 2):

$$\begin{array}{r} 01101001 \text{ ROL } 2 \\ \hline 10100101 \end{array}$$

- Preserves all bits, useful in cryptography



# Rotate Right

- Shifts all bits to the right, wrapping around
- Example (8-bit, rotate by 2):

$$\begin{array}{r} 01101001 \text{ ROR } 2 \\ \hline 10100101 \end{array}$$

- Example (8-bit, rotate by 2):
- Also preserves all bits, used in various algorithms



# Arithmetic Left Shift

- Shifts bits to left just like logical shift
- If sign changes, it's an overflow
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \lll 1 \\ \hline 11010010 \end{array}$$

- Equivalent to multiplication by  $2^n$  for  $n$  shifts
- Can cause overflow if significant bits are shifted out





# Arithmetic Right Shift

- Shift bits right preserving sign
- Leftmost bits are filled with the sign bit
- Example (8-bit, shift by 2):

$$\begin{array}{r} 11101001 \ggg 1 \\ \hline 11110100 \end{array}$$

- For positive numbers, equivalent to division by  $2^n$  for  $n$  shifts
- For negative numbers, rounds towards negative infinity
- Preserves the sign of the number



# Bitwise Operations in Verilog

- Verilog provides operators for various bitwise operations
- These operations are fundamental in digital design
- Examples (assuming 8-bit variables a and b):

```
c = a & b;           // Bitwise AND
d = a | b;           // Bitwise OR
e = a ^ b;           // Bitwise XOR
f = ~a;              // Bitwise NOT
g = a << 2;           // Logical left shift
h = a >> 2;           // Logical right shift
i = $signed(a) >>> 2; // Arithmetic right shift
j = {a[5:0], a[7:6]}; // Rotate left
k = {a[1:0], a[7:2]}; // Rotate right
```



- ASCII: American Standard Code for Information Interchange
- Maps characters to binary values
- ASCII table overview <https://www.ascii-code.com/>
- Example: 'A' = 01000001 = 65
- Example: 'B' = 01000010 = 66
- Example: 'a' = 01100001 = 97



# Unicode UTF-8 Encoding

- Unicode is a standard for representing text of all languages
- Originally fit into 16 bits
- Now requires 19 bits because of emoji
- UTF-8 is a variable length encoding
- Useful when most of the text is English
- <https://symbl.cc/en/unicode-table/>
- <https://r12a.github.io/app-conversion/>

1st byte	2nd byte	3rd byte	4th byte	19-bit value
0xxxxxxx				0000000000000000xxxxxx
110yyyyy	10xxxxxx			0000000000yyyyyxxxxxx
1110zzzz	10yyzzzz	10yyyyyy	10xxxxxx	00000zzzzyyyyyyxxxxxx
11110uuu	10uuzzzz	10yyyyyy	10xxxxxx	uuuuuzzzzyyyyyyxxxxxx



# Unicode UTF-8 Encoding

Let's decode the following bytes

00000000: 74 65 73 74 0a ce b5 cf 85 cf 87 ce b1 cf 81 ce b

00000010: cf 83 cf 84 cf 8e 0a E5 A4 A7

- 74 = 01110100 starts with 0, so look up ASCII
- same through 0A (ASCII for linefeed)
- CE B5 = **11**001110 **10**110101 starts with 110, so it's a 2-byte char
- bits = 01110 110101 = 011 1011 0101 = 3B1
- Look up 3B1 in Unicode online (it's a Greek letter)
- E5 A4 A7 = **11**100101 **10**100100 **10**100111 starts with 1110, so it's a 3-byte char
- bits = 0101 100100 100111 = 5 9 2 7
- Look up 0x5927 in Unicode online (Chinese)



# Fixed Point Fractions

$$2^3 \mid 2^2 \mid 2^1 \mid 2^0 \mid . \mid 2^{-1} \mid 2^{-2}$$

- Fractions in binary are represented as negative powers
- $\frac{1}{2}, \frac{1}{4}, \frac{1}{8}$
- Examples

$$101.1 = 4 + 1 + \frac{1}{2} = 5.5$$

$$1.01 = 1 + \frac{1}{4} = 1.25$$

$$110.11 = 4 + 2 + \frac{1}{2} + \frac{1}{4}$$

$$1001.001 = 8 + 1 + \frac{1}{8}$$



# Floating Point

- Fixed point represents fractions, but only a single size
- Floating point can represent values wildly different values
- IEEE-754

<https://www.h-schmidt.net/FloatConverter/IEEE754.html>

Single precision	32 bits
Double precision	64 bits
Quad Precision	128 bits (not yet in hardware)
Half Precision	16 bits (GPUs)
fp8	8 bits (GPUs)
fp4	4 bits (GPUs)



# R

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻ 40/43



- Infinity is a special value
- Too large to be countable
- From calculus:  $\lim_{x \rightarrow 0} \frac{1}{x} = \infty$
- Example:  $1.0/0.0 = \infty$
- Not countable:  $1.0/0.0 = 2.0/0.0$
- Negative infinity:  $-1.0/0.0 = -\infty$



# NaN: Not a Number

- NaN is a special value meaning the answer is unknown
- When two opposing infinities fight, we don't know the answer
- Example  $1.0/0.0 - 1.0/0.0 = \text{NaN}$
- Example:  $0.0/0.0 = \text{NaN}$



# NaN: Not a Number

For each of the following, determine if the result is NaN, Infinity, or a number

$x = 5.0/0.0$	
$y = 1.0/x$	
$z = \sqrt{x}$	
$w = \sin(x)$	
$v = -3.5/0.0$	
$u = \sqrt{v}$	

